Lumi · continuation plan

Lumi already runs end to end on one machine: it captures fixes from Slack, Teams, e-mail and pasted threads, links them to the equipment and people involved, checks them with one Yes/No question at a time, and serves the confirmed knowledge to colleagues and to AI agents over MCP. What it has not yet shown is the part that matters most: real engineers answering over weeks, on a live channel. The next step is an eight-week pilot with one Elisa operations team of about ten volunteers. Lumi runs inside Elisa's environment, only reads one channel and one mailbox label, asks each person at most a few questions a day and never at night, and can be switched off with one click.

During the pilot we measure how often engineers answer and how many seconds it costs them, how many confirmed fixes Lumi keeps each week, how often a person or an agent reuses one, and every case Lumi links or confirms wrongly, reviewed weekly with the team lead. In parallel we improve the weakest step, deciding whether an old fix applies to a new unit, and learn which questions to ask from the real answers, keeping a change only if it beats the current rules on cases it has not seen. If the numbers hold, Lumi extends to more teams and channels, gains single sign-on and per-team permissions, and becomes the source of confirmed operational knowledge for Elisa's own agents.