Lumi · continuation plan

Lumi already runs end to end on one machine: it captures fixes from Slack, Teams, e-mail and pasted threads, links them to the equipment and people involved, checks them with one Yes/No question at a time, and serves the confirmed knowledge to colleagues and to AI agents over MCP. A next step could be running it with a real team on a live channel, measuring how often engineers answer, how much of their time it takes and how often a confirmed fix gets reused. From there it could improve its weakest step, deciding whether an old fix applies to a new unit, learn which questions to ask from real answers, and grow to more teams and channels, becoming the source of confirmed knowledge for the company's own AI agents.